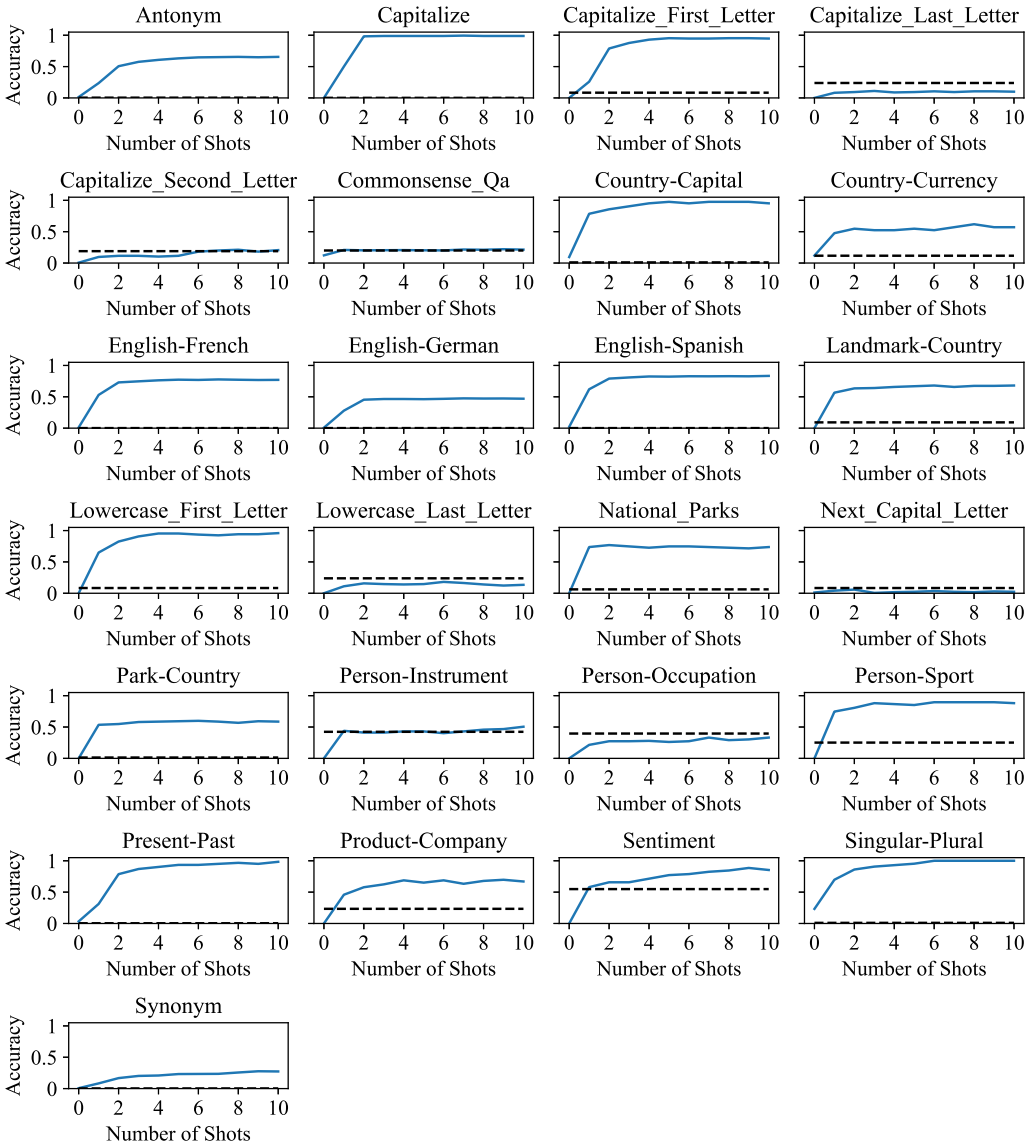

Figure 8: Few-shot ICL performance (top-1 accuracy) for GPT-J on a set of 25 abstractive-style tasks. In general, more shots improves performance. However, for many of these tasks the accuracy plateaus after a number of shots. The dotted baseline shows the accuracy of predicting only the majority label.

E.2 Few-Shot ICL Performance

Figure 8 shows the few-shot performance (top-1 accuracy) for GPT-J on a larger subset of our task list. The dotted baseline is based on the majority label for the dataset, computed as (# majority label)/(# total instances).

Figure 9 shows the few-shot performance (top-1 accuracy) for GPT-J on additional extractive tasks. For datasets with lists of words as input, the dotted baseline is computed to be $1/\text{size of the input list}$, (e.g. $1/3, 1/5$), and for all other tasks it represents predicting the majority label of the dataset, computed as (# majority label)/(# total instances).

E.3 Evaluating Function Vectors on Additional Tasks

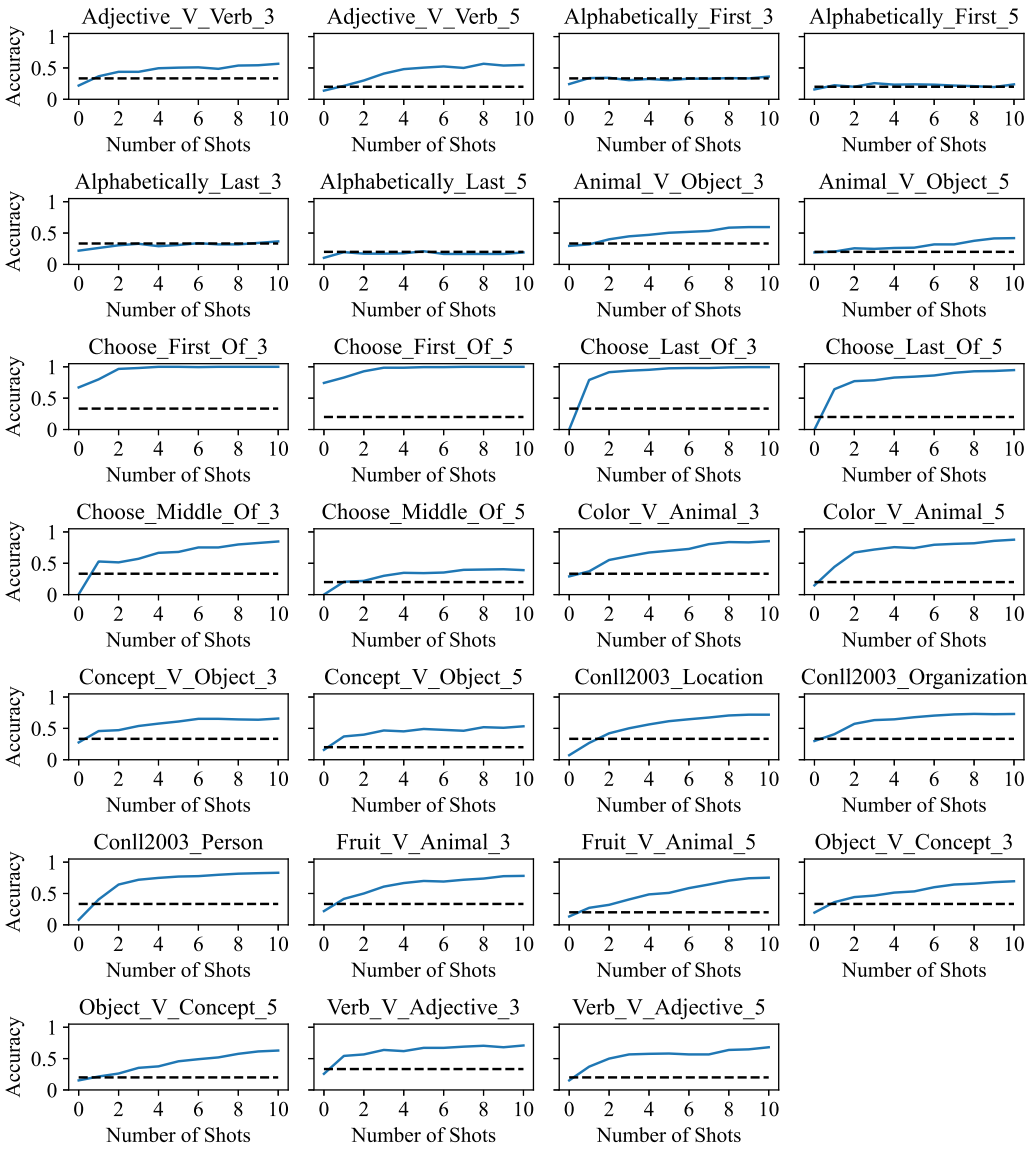

Figure 9: Few-shot ICL performance (top-1 accuracy) for GPT-J on a set of 27 extractive-style tasks. In general, more shots improves model performance. A dataset with 3 or 5 at the end denotes the size of the input list of words used. There are a few cases where the model cannot perform the task any better than random (e.g. choose the alphabetically first word in a list), which we do not analyze further. The dotted baseline shows the accuracy $1/(\text{list size})$, or the majority label in the case of the datasets derived from conll2003.

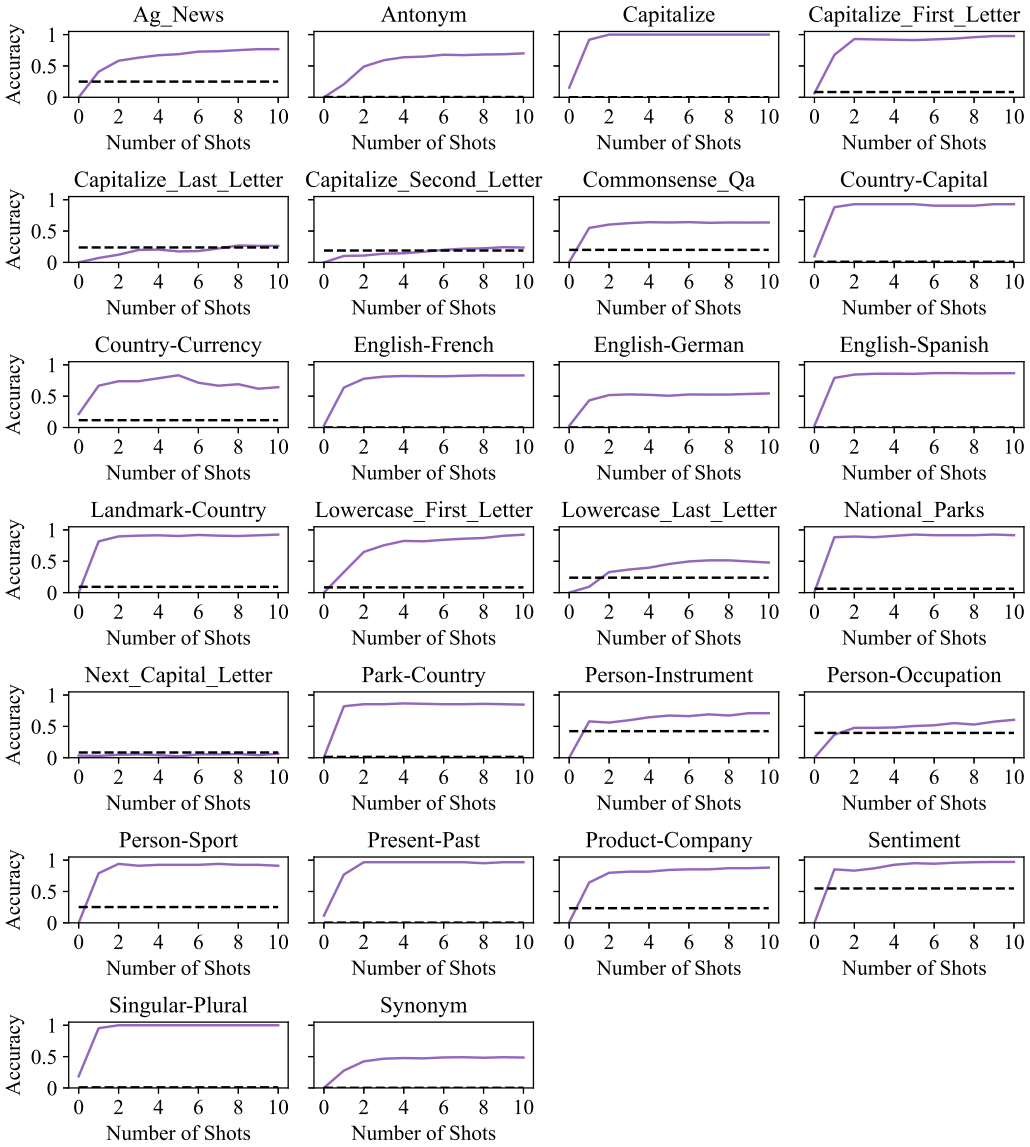

Figure 10: Few-shot ICL performance (top-1 accuracy) for Llama 2 (13B) on a set of 26 abstractive-style tasks. Using more ICL examples (shots) improves performance for many of these tasks, though the performance does plateau after a few shots for many of the tasks. There are a few cases where Llama 2 (13B) cannot perform the task any better than random (e.g. capitalize the second letter in the word), which we do not analyze further. The dotted baseline shows accuracy choosing the majority label.