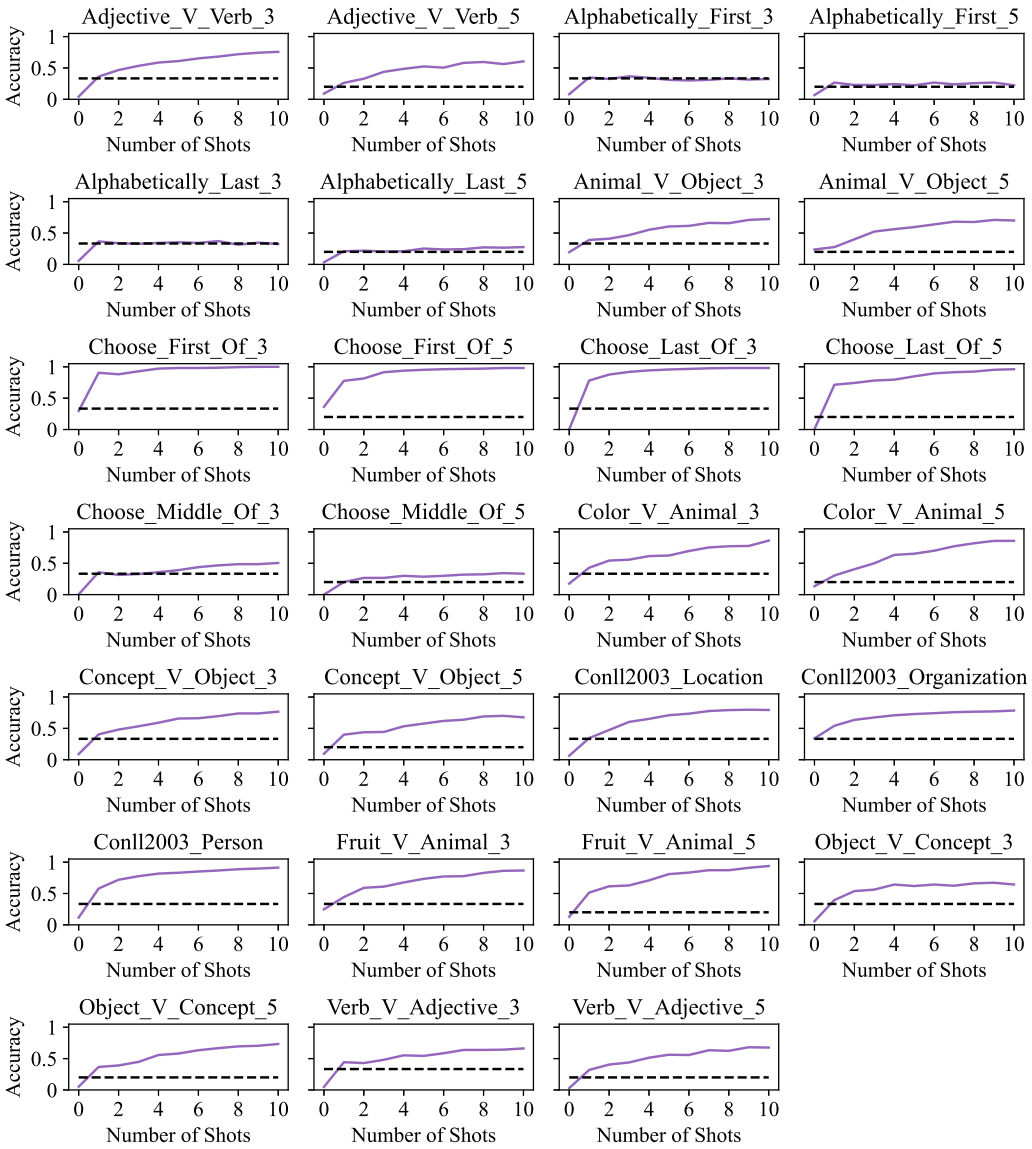

Figure 11: Few-shot ICL performance (top-1 accuracy) for Llama 2 (13B) on a set of 27 extractive-style tasks. Using more ICL examples improves performance for most of the tasks. A dataset name ending with 3 or 5 denotes the size of the input word list. There are a few cases where the model cannot perform the task any better than random (e.g. choose the alphabetically first word in a list), which we do not analyze further. The dotted baseline shows the accuracy $1/(\text{list size})$ (e.g. $1/3, 1/5$, or the majority label in the case of the conll2003 datasets.

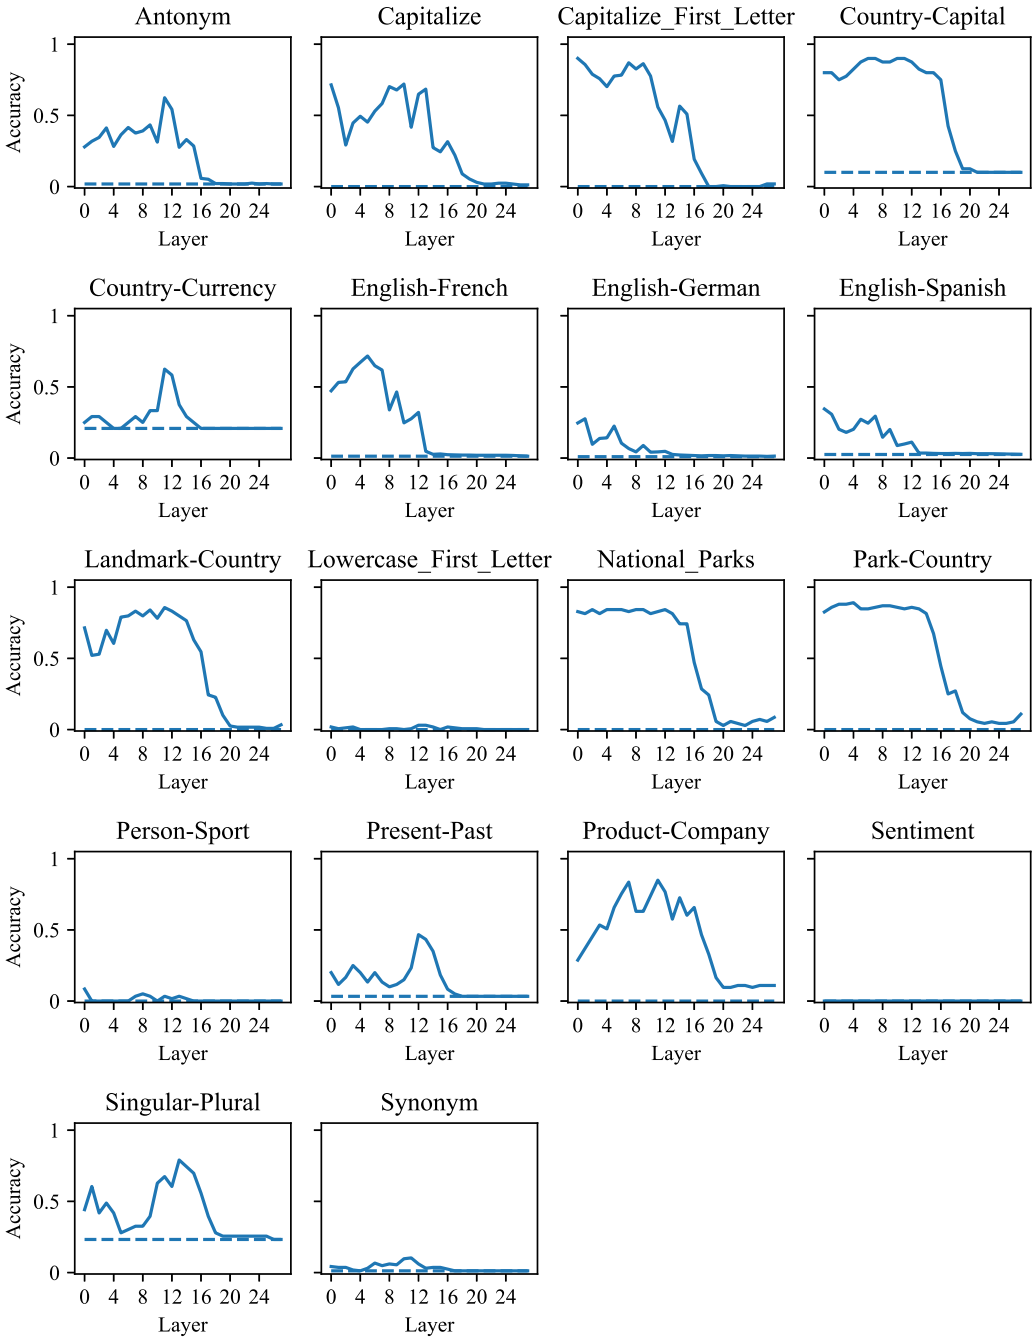

Figure 12: Zero-Shot Top-1 Accuracy Results of adding FVs to GPT-J across 18 of our task datasets. In addition the 6 analyzed in the main paper, we see similar results and performance across a variety of other tasks - adding FVs in early-middle layers seems to have the most effect. Notable exceptions include lowercase first letter, person-sport, synonym, and sentiment, where the FV doesn't seem to have much effect. Note we did not evaluate the FV on tasks where the model's ICL performance was poor (compared to the majority label baseline) (Figure 8).

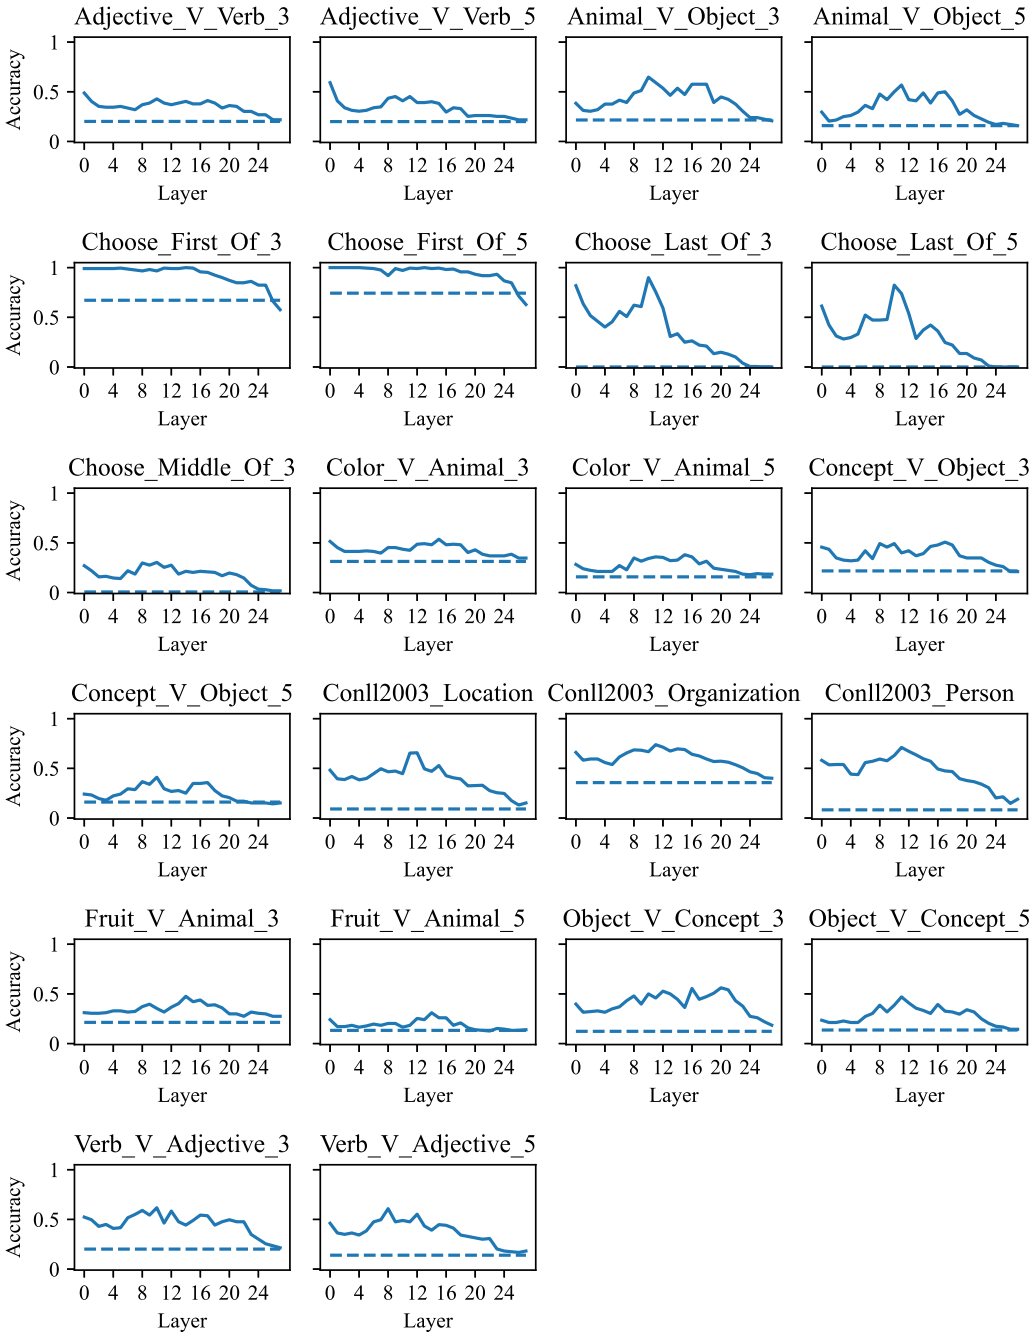

Figure 13: Zero-Shot Top-1 Accuracy Results of adding FVs to GPT-J across 22 of our task datasets. Shown here are mainly extractive tasks, and we see similar trends across all tasks. The FV performance is fairly stagnant across layers, with many tasks having peak performance in middle layers of the network. The model's zero-shot performance without intervention is plotted as a dotted line. Adding the FV causes the model to extract the correct entity much more often than the base model in this uninformative zero-shot case.